

Context Memory Dashboard

Author: Paridhi Talwar, Product Manager

Status: In Review

Last Updated: July 13, 2026

Target Release: v1 – Q3 2026

1. Problem Statement

Every AI assistant that remembers things across conversations is sitting on the same unsolved problem: the memory is invisible. The model gets smarter about you with every chat, but you have no window into what it actually knows, and no way to fix it when it's wrong.

Right now, memory happens to the user. It doesn't happen with them.

- **Current Behavior:** A user mentions they're job hunting in one chat → Three weeks later, in a completely unrelated conversation, the assistant references it unprompted → The user is caught off guard.
- **The Gap:** There's no screen anywhere that shows what's been remembered. If the assistant gets something wrong — an old job, a preference that's changed, a health detail mentioned once and never again — there's no direct way to correct it. The only lever most users have is a single, blunt "turn memory off" switch that throws away everything, good and bad.
- **What we're hearing:** In conversations with users who've used memory-enabled assistants for a while, the same two lines come up over and over: "I didn't know it remembered that," and "how do I make it forget this one thing." Nobody's asking for less memory. They're asking for a way to see it.

The missing piece isn't a smarter memory. It's a mirror. This proposal is about building that window.

Memory Dashboard turns the assistant's background memory into something the user can actually see and edit, a settings surface, not a new AI feature. No new model capability required, no retraining. Just visibility into something that already exists.

The business case: right now we can tell users "the assistant remembers you." With this, we can tell them exactly what it remembers, and prove we'll let them control it. That's the difference between a feature people tolerate and one they trust enough to put sensitive context into.

2. Target Audience & Segmentation

Segment	Description	Primary Need
Power users	Daily / multi-session users who rely on continuity (e.g. ongoing projects, coaching, tutoring use cases)	Confidence that useful context persists correctly
Privacy-conscious users	Users discussing sensitive topics (health, legal, career, finances)	Visibility and granular control / deletion

Segment	Description	Primary Need
New / occasional users	Infrequent users who notice memory as a "surprise" feature	A simple explanation of what memory is and how to turn it off

We're prioritizing power users and the privacy-conscious for v1 because they are the ones actively asking for this control. For casual users, a full dashboard is overkill—a simple chat tooltip is enough to keep them informed.

3. User's Current State Journey Map

Stage	Current Experience	Pain Point
Discovery	User notices assistant referencing something from a past session	No warning, no way to verify accuracy
Reaction	User is either pleased or unsettled	No path to act on either reaction
Attempted control	User digs through generic settings	Memory controls are buried, coarse (all-or-nothing), or absent
Resolution	User either accepts it or disables memory entirely	Disabling loses all value, including useful context

4. Proposed Solution

A Context Memory Dashboard, accessible from account settings and from an in-conversation shortcut, with three core capabilities:

4.1 View

- A categorized, human-readable UI (e.g., Work, Preferences, Code) accessible from settings.
- Technical note: We will not expose raw embeddings or logs. The backend will need to translate the vector data (e.g., ChromaDB entries) into natural language summaries for the frontend.

4.2 Edit & correct

- Inline edit or delete of any individual memory item, with immediate effect on future conversations.
- A lightweight "this is wrong" flag that lets users flag inaccurate memories without needing to know the exact correct phrasing.

4.3 Control & scope

- Per-category toggles (e.g. turn off retention for "Health" topics specifically, keep it on for "Work").
- A visible "Incognito" or no-memory mode indicator so users always know their current memory state.
- Bulk export and bulk delete, to satisfy data-portability and right-to-erasure expectations.

4.4 Explicit non-goals for v1

- No user-facing exposure of the underlying retrieval/ranking mechanics (i.e. this is not a debugging tool).

- No cross-user or team-shared memory management — single-user scope only.
- No real-time memory editing from within an active conversation (that is a v2 consideration, noted in Trade-offs).

5. Wireframe & User Flow

Primary flow: Settings entry point

- Settings → "What I remember" → Dashboard home (categorized list, search bar, most-recent-first)
- Dashboard home → tap a memory item → Detail view (full text, source conversation link if available, Edit / Delete / Flag actions)
- Dashboard home → category toggle → Confirmation modal ("Turning this off will stop new [Health] memories from being saved. Existing ones remain until deleted.") → Toggle applied
- Dashboard home → "Delete all" → Two-step confirmation (type to confirm) → All memories cleared, dashboard shows empty state with explanation

Secondary flow: In-conversation entry point

- User taps a small "memory" icon shown when the assistant references a past fact → Slide-over panel shows the specific fact used, with a direct link to edit/delete that item → returns to conversation

(Note: full annotated wireframes for these flows are maintained as a companion artifact — see Wireframes project — since visual layout is easier to review as images than as a text description.)

6. Success Metrics

Metric	Definition
Dashboard adoption	% of active users who open the dashboard at least once
Correction usage	% of dashboard visitors who edit, delete, or flag at least one item
Trust signal	Change in survey-measured trust in memory accuracy, pre vs. post
Support deflection	Reduction in support tickets tagged "memory / personalization confusion"
Retention guardrail	Memory-enabled user retention, pre vs. post (must not regress)
Full opt-out rate	% of users who disable memory entirely after visiting dashboard

Guardrail metric: latency and load time of the dashboard itself should not exceed 1.5s at p95 — a slow "trust" surface undermines the feature it's meant to support.

7. Launch Readiness

7.1 Phased rollout

Phase	Audience	Exit Criteria
Alpha	Internal team + 50 opted-in power users	No P0/P1 bugs; edit/delete actions correctly reflected in next conversation 100% of time in testing
Beta	5% of memory-enabled users, randomized	Dashboard adoption $\geq 10\%$; no regression in overall retention; support ticket volume stable or down
GA	100% of memory-enabled users	Beta success metrics hold for 2 consecutive weeks

7.2 Cross-functional readiness checklist

- **Engineering:** Ensure edit and delete actions are immediate. Data needs to be wiped from the retrieval layer entirely, not just hidden from the user's view.
- **Legal & Privacy:** Review our language and data promises. We need to be crystal clear about what "delete" actually means so users can trust our commitments.
- **Design:** Map out the tough moments, not just the easy ones. We need clear plans for empty states, errors, sync delays, and confirmation screens.
- **Support:** Get ahead of the questions. Prepare FAQs and scripts for users wondering why memory didn't update or why old facts are sticking around.
- **Analytics:** Track everything. Ensure every action defined in our success metrics is working perfectly in staging before we hit the Alpha phase.

8. Trade-offs & Open Questions

8.1 Key trade-offs

- **Control vs. Simplicity:** Giving users control over every single fact is powerful, but it's too much for most people. For v1, we're sticking to category-level controls. We'll make individual edits available but won't put them front and center.
- **Instant vs. Delayed Sync:** It would be great if edits happened instantly during a chat, but that's technically risky and messy. We're going with a slight delay—edits will apply to the *next* conversation—to keep things reliable and smooth.
- **Transparency vs. Complexity:** We want users to trust the system without needing a degree in computer science. Showing them simple summaries is the right balance. We're accepting the risk that these summaries might not perfectly match the "raw" data the system retrieves; we'll watch that gap closely after we launch.

8.2 Open questions for review

- **Handling corrections:** If a user flags a memory as "wrong," should it vanish immediately, or wait for a check to confirm it?
- **Onboarding:** Should new users get the dashboard on day one, or should we wait until they've actually built up enough memory to care about it?
- **Enterprise needs:** Do businesses need an admin-controlled version, or is this strictly a single-user tool until we see how people actually use it?